

AI Level Director Studio: A Designer Governed Workflow for 2D Platformer Level Iteration

Student: Robert Mayfield

Industry Context and Problem Definition

AI Level Director Studio targets independent and small team game development, a sector under pressures that make responsible AI integration urgent. In the Game Developers Conference 2025 State of the Game Industry survey of more than 3,000 developers, 11% had been laid off in the past year and 52% worked at companies using generative AI. Adoption is rising but confidence is not: 30% believed generative AI was having a negative impact, up 12 points year over year, citing intellectual property theft, falling quality, and bias (Game Developers Conference, 2025).

That combination, more generated content, fewer people, and falling trust, is the real problem to design for. Level design is part science, part craft: difficulty, pacing, and a safe opening can be measured, but whether a segment is genuinely engaging is a creative judgment built on a designer’s experience. Mixed initiative tools have long treated level design as a collaboration between designer and system rather than full automation (Smith et al., 2010). Review and playtesting are expensive and leave a small team little slack. The hard problem is not generation: raw generated content cannot be trusted as a finished asset, because it may be structurally invalid, misaligned with intent, or derivative of its training data.

The need is therefore the move from “an AI can produce a candidate” to “a designer can responsibly decide what to do with it.” That is an applied system design problem under real constraints: a small team needs something lightweight and transparent, and originality and intellectual property risk must be disclosed rather than hidden. AI Level Director Studio helps a designer screen, triage, and iterate on candidate segments faster while keeping every creative decision in human hands.

Overview of the Integrated Solution

AI Level Director Studio is a designer governed workflow for iterating on 2D platformer level segments. Rather than a new model, it is a compound AI system, reflecting the observation that many effective AI applications combine models, tools, and orchestration rather than one monolithic model (Zaharia et al., 2024). It contributes the orchestration layer that turns three prior AI projects into one coherent loop, plus the candidate lifecycle, persistence, reporting, and interface that make the loop usable and transparent.

The workflow runs from a brief to generated or uploaded candidates, agentic triage, a playtest decision, feedback classification, and finally completing or revising the candidate. Every candidate carries its own lifecycle state, so several can sit at different stages within one session.

The system is deliberately advisory. It recommends and routes, but never edits a level, never finalizes one, and never claims to predict whether a level is fun. Its outputs are observable: a candidate board, JSON session snapshots, an append only event log, and a Markdown report. The same backend drives a reproducible notebook and a Gradio interface (Figure 1).

Integration of Prior Projects and Methods

The system integrates three of my prior capstone projects, each wrapped behind an adapter so the studio coordinates them without rewriting them. In each case a lesson from the original project shaped its role here.

1. **Generative AI, the candidate source.** The conditional Transformer level generator proposes candidate segments. Procedural content generation via machine learning serves not only generation but co-creation, critique, and analysis, and is constrained by scarce training data for custom game content (Summerville et al., 2018). My results showed this: the generator produced structurally valid, difficulty controlled segments, but novelty analysis revealed substantial similarity to the training levels. The studio therefore treats generated levels as drafts and discloses their derivative risk.




Figure 1: System architecture: the UI and notebook call one service facade, which reaches the prior projects through adapters and persists each step.

2. **Agentic AI, the triage authority.** The level design triage agent is the design judgment layer and the single triage authority. It interprets a brief, gathers facts with deterministic tools, reasons about how a candidate serves the intent, and recommends one of six actions, interleaving reasoning with tool grounded actions rather than relying on internal knowledge alone (Yao et al., 2022). The lesson I carried forward is that the agent’s value is judgment, not facts, so it is reused unchanged and never duplicated.
3. **Machine Learning, the feedback signal.** The player feedback classifier supplies the post playtest reception signal, classifying playtester text as positive or negative. Because a sentiment classifier is useful but narrow, it runs only after a candidate is sent to playtest and is treated as a broad signal, not a complete playtest analysis.

In sequence, Generative AI proposes, Agentic AI judges, and Machine Learning reports back, while the studio owns the lifecycle, persistence, and reporting that connect them.

Technical Design Decisions and Tradeoffs

The decisions that shaped the system were about its overall design and what to integrate, each weighing a real alternative.

Decision A: A designer governed loop, not an autonomous engine. A pipeline that generates and finalizes levels on its own was rejected because level design is part craft and the designer must keep creative authority. The system advises and routes; it never decides. This is the spine of the synthesis.

Decision B: Treat every signal as evidence, never a verdict. The alternative, a single score that approves candidates, fails because no model can judge whether a level is fun, and a small studio has little reason to trust a black box that claims to. Validity, difficulty, novelty, pacing, and sentiment are evidence for the designer to weigh, consistent with human AI interaction guidelines that stress user control and clear limits (Amershi et al., 2019).

Decision C: A designer facing UI, not a notebook only demonstration. A notebook would prove the pipeline runs, but a UI makes the independent but connected nature of the components tangible: a designer drives the loop and sees each system’s contribution, which builds trust in an AI assisted tool.

Decision D: Non-destructive, fully audited iteration. Rather than one mutable level, each candidate owns its state, revisions create new linked candidates, and every action appends to an event log (Figure 2). Overwriting as you go is simpler but loses the variant history designers iterate on and the provenance accountability needs.

Decision E: Composable, swappable components behind adapters. Wiring the three systems together directly invites the entanglement that makes machine learning systems fragile (Sculley et al., 2015). Adapters let a small team upgrade or replace one model without rebuilding the system, keeping the orchestration dependent on stable interfaces, not model internals.




Figure 2: Candidate lifecycle state machine. Triage and feedback drive transitions; the designer can complete or archive from any non terminal state.

Evaluation and Reflection

Across seven scenarios (A, A2, B through F) the loop behaved as designed: a generated candidate was rejected for a structural flaw before playtest, a candidate that duplicated a known level was caught as derivative and routed back, and an ambiguous brief drew a clarification request rather than a forced recommendation. On the feedback path, positive feedback completed a candidate and negative feedback created a linked revision; when the classifier misread a short note as positive, the designer overrode it and its original label was kept on the record. The designer governed loop held, routing evidence to the human and never finalizing a level itself, which shows sensible behavior, not improved design outcomes.

The harder work was in the seams, not the models. The three systems shared no common notion of a candidate, and each returned a different output, raw tile text, a structured action, a bare sentiment label, with different certainty semantics. Reconciling those into one advisory lifecycle, and routing weak or conflicting signals without letting any dominate, was harder than running any single model. The lesson: integrating across domains pays off in the orchestration and the boundaries, not the modeling.

Ethical, Governance, and Responsible AI Considerations

Derivative content is the clearest risk. The generator was trained on a small corpus and its outputs showed high similarity to it, so a generated candidate can closely resemble an existing level, an intellectual property risk developers themselves flagged (Game Developers Conference, 2025). The studio addresses it directly: generated levels are labeled drafts, their derivative risk is surfaced as a warning rather than hidden, and a too similar candidate can be routed to human originality review instead of accepted automatically.

Over trust, treating a model’s output as a verdict, is the second risk. The feedback classifier is a broad reception signal, not a complete analysis, and the triage agent offers judgment, not ground truth about fun. In line with human AI interaction guidelines on transparency, user control, and clear limits (Amershi et al., 2019), the system never finalizes a level, supports clarification and human review states, saves each triage’s full reasoning transcript for audit, and discloses its limitations in every report (Mitchell et al., 2019). The classifier is the clearest case: because my own analysis found it misreads short reviews, short feedback is flagged and the designer can override the classified sentiment after reading the text, with the model’s original label kept on the record.

Responsible deployment requires more than the prototype provides: broader training data, far more playtesting, asset license review, privacy handling for real feedback, and professional quality assurance. Accountability stays with the human designer.

Limitations and Risks

This is a prototype, and its limits are real.

- **Generator:** trained on a small corpus, so its candidates carry derivative risk and are not production ready content.
- **Classifier:** a bag of words model trained on consumer game reviews, so it misses negation and sarcasm, shifts domain when applied to playtest notes, and misreads roughly 30% of short negative reviews as positive. It returns a label only, with no calibrated confidence.
- **Triage agent:** its difficulty estimate is heuristic and structural, not player validated, and because it uses a language model its judgments vary between runs and depend on an external provider.
- **Playtesting:** simulated rather than connected to real players.

The main standing risk is over reliance: if a team treated these advisory signals as final decisions, the system's value would invert. The designer governed design mitigates that risk but does not remove it.

Professional and Industry Relevance

The pressing question in game development today is not whether AI can generate content, it clearly can, but whether a studio can adopt it without sacrificing originality, quality, or team trust. AI Level Director Studio is a small, working answer. It treats AI as designer governed assistance that accelerates iteration while keeping creative authority, transparency, and accountability with the human: generated content is screened, not shipped; every signal is evidence, not a verdict; and the full decision history is recorded for review. Integrating specialized models into an accountable workflow, rather than a system that designs on its own, is what studios increasingly need, and the pattern generalizes to any creative or high stakes field adopting AI responsibly.

Future Extensions or Improvements

Several improvements follow from the limitations above.

1. **Calibrated confidence on the classifier**, for example through Platt scaling, so it could express how certain a feedback judgment is and route uncertain cases to human review.
2. **A stronger generator** that reduces derivative risk, for instance a larger pretrained model rather than a small from scratch one.
3. **Live telemetry from the deep learning project** to complement post playtest sentiment, and connecting the playtester step to real players for proper user studies of whether the workflow improves design outcomes.
4. **Lighter polish:** image or sprite based previews using appropriately licensed assets, and an optional report narrator that narrates around the ground truth facts rather than generating them.

References

- Amershi, S., Weld, D., Vorvoreanu, M., Fourney, A., Nushi, B., Collisson, P., Suh, J., Iqbal, S., Bennett, P. N., Inkpen, K., Teevan, J., Kikin-Gil, R., & Horvitz, E. (2019). Guidelines for human-AI interaction. *Proceedings of the 2019 CHI Conference on Human Factors in Computing Systems*, 1-13. <https://doi.org/10.1145/3290605.3300233>
- Game Developers Conference. (2025). *GDC 2025 state of the game industry: Devs weigh in on layoffs, AI, and more*. <https://gdconf.com/article/gdc-2025-state-of-the-game-industry-devs-weigh-in-on-layoffs-ai-and-more/>
- Mitchell, M., Wu, S., Zaldivar, A., Barnes, P., Vasserman, L., Hutchinson, B., Spitzer, E., Raji, I. D., & Gebru, T. (2019). Model cards for model reporting. *Proceedings of the Conference on Fairness, Accountability, and Transparency*, 220-229. <https://doi.org/10.1145/3287560.3287596>
- Sculley, D., Holt, G., Golovin, D., Davydov, E., Phillips, T., Ebner, D., Chaudhary, V., Young, M., Crespo, J.-F., & Dennison, D. (2015). Hidden technical debt in machine learning systems. *Advances in Neural Information Processing Systems*, 28. <https://papers.nips.cc/paper/5656-hidden-technical-debt-in-machine-learning-systems.pdf>
- Smith, G., Whitehead, J., & Mateas, M. (2010). Tanagra: A mixed-initiative level design tool. *Proceedings of the Fifth International Conference on the Foundations of Digital Games*, 209-216. <https://doi.org/10.1145/1822348.1822376>
- Summerville, A., Snodgrass, S., Guzdial, M., Holmgård, C., Hoover, A. K., Isaksen, A., Nealen, A., & Togelius, J. (2018). Procedural content generation via machine learning (PCGML). *IEEE Transactions on Games*, 10(3), 257-270. <https://doi.org/10.1109/TG.2018.2846639>
- Yao, S., Zhao, J., Yu, D., Du, N., Shafran, I., Narasimhan, K., & Cao, Y. (2022). *ReAct: Synergizing reasoning and acting in language models*. <https://arxiv.org/abs/2210.03629>
- Zaharia, M., Khattab, O., Chen, L., Davis, J. Q., Miller, H., Potts, C., Zou, J., Carbin, M., Frankle, J., Rao, N., & Ghodsi, A. (2024). *The shift from models to compound AI systems*. Berkeley Artificial Intelligence Research. <https://bair.berkeley.edu/blog/2024/02/18/compound-ai-systems/>